

Divya Nallawar

Student ID: 110988624

Application of Statistical MOdelling Techniques on a Fashion related Dataset

This project focuses on understanding and applying different statistical models on a dataset. The dataset that I am using is obtained from Kaggle. This dataset was collected simply to learn the different trends in the fashion.

```
# Loading all the packages that would be used over the project
library(tidyverse)
library(caret)
library(corrplot)
library(readr)
library(dplyr)
library(tidyr)
library(lmtest)
library(car)
library(leaps)
library(MASS)

— Attaching core tidyverse packages
tidyverse 2.0.0 —
✓ dplyr     1.1.4    ✓ readr     2.1.5
✓forcats   1.0.0    ✓ stringr  1.5.1
✓ ggplot2   3.5.1    ✓ tibble    3.2.1
✓ lubridate 1.9.3    ✓ tidyrr    1.3.1
✓ purrr    1.0.2
— Conflicts

tidyverse_conflicts() —
✖ dplyr::filter() masks stats::filter()
✖ dplyr::lag()   masks stats::lag()
ⓘ Use the conflicted package (<http://conflicted.r-lib.org/>) to
force all conflicts to become errors
Loading required package: lattice

Attaching package: 'caret'

The following object is masked from 'package:purrr':
```

```
lift

corrplot 0.92 loaded

Loading required package: zoo

Attaching package: 'zoo'

The following objects are masked from 'package:base':
  as.Date, as.Date.numeric

Loading required package: carData

Attaching package: 'car'

The following object is masked from 'package:dplyr':
  recode

The following object is masked from 'package:purrr':
  some

Attaching package: 'MASS'

The following object is masked from 'package:dplyr':
  select

# Here to fetch the data, one more possible way is to use the Kaggle API to fetch the data. But here, we have downloaded the file directly from kaggle and using it.
# Read the downloaded CSV file
myntra_fashion_dataset <- read_csv("./Myntra_Fasion_Clothing.csv",
  col_types = cols(DiscountOffer = col_character()))
print(head(myntra_fashion_dataset))

# A tibble: 6 × 13
  URL      Product_id BrandName Category Individual_category
```

```

category_by_Gender
<chr>          <dbl> <chr>      <chr>      <chr>      <chr>
1 https://...    2296012 Roadster  Bottom ... jeans        Men
2 https://...    13780156 LOCOMOTI... Bottom ... track-pants   Men
3 https://...    11895958 Roadster  Topwear  shirts       Men
4 https://...    4335679 Zivame    Lingeri... shapewear   Women
5 https://...    11690882 Roadster  Western  tshirts     Women
6 https://...    2490950 Mast & H... Western  tops        Women

# ⓘ 7 more variables: Description <chr>, `DiscountPrice (in Rs)` <dbl>,
#   `OriginalPrice (in Rs)` <dbl>, DiscountOffer <chr>, SizeOption <chr>,
#   Ratings <dbl>, Reviews <dbl>

colnames(myntra_fashion_dataset) # Fetching the column name from the dataset

[1] "URL"                  "Product_id"           "BrandName"
[4] "Category"              "Individual_category"
"category_by_Gender"
[7] "Description"           "DiscountPrice (in Rs)" "OriginalPrice
(in Rs)"
[10] "DiscountOffer"         "SizeOption"            "Ratings"
[13] "Reviews"

summary(myntra_fashion_dataset) # let's look at the summary of the dataset before performing any datacleaning


```

URL	Product_id	BrandName	Category
Length:526564	Min. : 27399	Length:526564	
Length:526564	1st Qu.:13880530	Class :character	
Class :character	Median :15971057	Class :character	
Mode :character	Mode :character	Mode :character	
Mode :character	Mean :15069387		
	3rd Qu.:17347414		
	Max. :18464352		

```

Individual_category category_by_Gender Description
Length:526564      Length:526564      Length:526564
Class :character    Class :character    Class :character
Mode  :character    Mode  :character    Mode  :character

DiscountPrice (in Rs) OriginalPrice (in Rs) DiscountOffer
Min.   : 127          Min.   : 99          Length:526564
1st Qu.: 659          1st Qu.: 1299        Class :character
Median  : 952          Median  : 1999        Mode  :character
Mean    : 1237         Mean    : 2414
3rd Qu.: 1469         3rd Qu.: 2899
Max.   :27996          Max.   :90000
NA's   :193158

SizeOption      Ratings      Reviews
Length:526564  Min.   :1.0  Min.   : 0
Class :character 1st Qu.:3.9  1st Qu.: 8
Mode  :character  Median :4.2  Median :18
                  Mean   :4.1  Mean   :62
                  3rd Qu.:4.4  3rd Qu.:52
                  Max.   :5.0  Max.   :999
                  NA's   :336152 NA's   :336152

missing_counts <- colSums(is.na(myntra_fashion_dataset))
print(missing_counts)

          URL          Product_id       BrandName
          0             0                 0
Category  Individual_category category_by_Gender
          0             0                 0
Description  DiscountPrice (in Rs) OriginalPrice (in Rs)
          0             193158            0
DiscountOffer      SizeOption      Ratings
          74306           0             336152
Reviews
          336152

```

Here, we can see that there are a lot of missing values in the dataset, specifically for columns namely `DiscountPrice (in Rs)`, `DiscountOffer`, `Ratings` and `Reviews`. Let's fix the missing values.

```

str(myntra_fashion_dataset)

spc_tbl_ [526,564 x 13] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
$ URL          : chr [1:526564]
"https://www.myntra.com/jeans/roadster/roadster-men-navy-blue-slim-
```

```

fit-mid-rise-clean-look-jeans/2296012/buy"
"https://www.myntra.com/track-pants/locomotive-men-black--white-solid-slim-fit-track-pants/13780156/buy"
"https://www.myntra.com/shirts/roadster/roadster-men-navy-white--black-geometric-printed-sustainable-casual-shirt/11895958/buy"
"https://www.myntra.com/shapewear/zivame/zivame-women-black-saree-shapewear-zi3023core0nude/4335679/buy" ...
$ Product_id : num [1:526564] 2296012 13780156 11895958
4335679 11690882 ...
$ BrandName : chr [1:526564] "Roadster" "LOCOMOTIVE"
"Roadster" "Zivame" ...
$ Category : chr [1:526564] "Bottom Wear" "Bottom Wear"
"Topwear" "Lingerie & Sleep Wear" ...
$ Individual_category : chr [1:526564] "jeans" "track-pants"
"shirts" "shapewear" ...
$ category_by_Gender : chr [1:526564] "Men" "Men" "Men" "Women" ...
$ Description : chr [1:526564] "roadster men navy blue slim
fit mid rise clean look jeans" "locomotive men black white solid slim
fit track pants" "roadster men navy white black geometric printed
sustainable casual shirt" "zivame women black saree shapewear
zi3023core0nude" ...
$ DiscountPrice (in Rs): num [1:526564] 824 517 629 893 NA NA 599 NA
NA NA ...
$ OriginalPrice (in Rs): num [1:526564] 1499 1149 1399 1295 599 ...
$ DiscountOffer : chr [1:526564] "45% OFF" "55% OFF" "55% OFF"
"31% OFF" ...
$ SizeOption : chr [1:526564] "28, 30, 32, 34, 36" "S, M,
L, XL" "38, 40, 42, 44, 46, 48" "S, M, L, XL, XXL" ...
$ Ratings : num [1:526564] 3.9 4 4.3 4.2 4.2 4.4 3.9 3.7
4.3 3.5 ...
$ Reviews : num [1:526564] 999 999 999 999 999 999 999 998
998 997 996 ...
- attr(*, "spec")=
  .. cols(
    .. URL = col_character(),
    .. Product_id = col_double(),
    .. BrandName = col_character(),
    .. Category = col_character(),
    .. Individual_category = col_character(),
    .. category_by_Gender = col_character(),
    .. Description = col_character(),
    .. `DiscountPrice (in Rs)` = col_double(),
    .. `OriginalPrice (in Rs)` = col_double(),
    .. DiscountOffer = col_character(),
    .. SizeOption = col_character(),
    .. Ratings = col_double(),
    .. Reviews = col_double()
  )
- attr(*, "problems")=<externalptr>

```

Part A : Data Cleaning

Here, we are cleaning our data, like removing the special characters from the dataset, filling the missing values, etc.

```
# Replace missing DiscountOffer values with 0
myntra_fashion_dataset$DiscountOffer <-
  ifelse(is.na(myntra_fashion_dataset$DiscountOffer), 0,
  myntra_fashion_dataset$DiscountOffer)

# Remove non-numeric characters from DiscountOffer and calculate
# discount percentage
myntra_fashion_dataset$DiscountOffer <- gsub("[^0-9]", "", 
myntra_fashion_dataset$DiscountOffer)
myntra_fashion_dataset$DiscountOffer <-
  as.numeric(myntra_fashion_dataset$DiscountOffer) /
  myntra_fashion_dataset$`OriginalPrice (in Rs)` * 100

# Handle missing values for 'DiscountPrice (in Rs)'
mask <- !is.na(myntra_fashion_dataset$`DiscountPrice (in Rs)`)) & !
  is.na(myntra_fashion_dataset$`OriginalPrice (in Rs)`)) &
  is.na(myntra_fashion_dataset$DiscountOffer)
myntra_fashion_dataset$DiscountOffer[mask] <-
  (myntra_fashion_dataset$`OriginalPrice (in Rs)`[mask] -
  myntra_fashion_dataset$`DiscountPrice (in Rs)`[mask]) /
  myntra_fashion_dataset$`OriginalPrice (in Rs)`[mask] * 100

# Calculate missing 'DiscountPrice (in Rs)' from 'DiscountOffer' and
# 'OriginalPrice (in Rs)'
mask <- !is.na(myntra_fashion_dataset$DiscountOffer) & !
  is.na(myntra_fashion_dataset$`OriginalPrice (in Rs)`)) &
  is.na(myntra_fashion_dataset$`DiscountPrice (in Rs)`))
myntra_fashion_dataset$`DiscountPrice (in Rs)`[mask] <-
  myntra_fashion_dataset$`OriginalPrice (in Rs)`[mask] * (1 -
  myntra_fashion_dataset$DiscountOffer[mask] / 100)

# Set missing values to 0 if both 'DiscountPrice (in Rs)' and
# 'DiscountOffer' are missing
mask <- is.na(myntra_fashion_dataset$`DiscountPrice (in Rs)`)) &
  is.na(myntra_fashion_dataset$DiscountOffer)
myntra_fashion_dataset$`DiscountPrice (in Rs)`[mask] <- 0
myntra_fashion_dataset$DiscountOffer[mask] <- 0

# Fill missing 'Ratings' and 'Reviews' with their mean values
myntra_fashion_dataset$Ratings <-
  ifelse(is.na(myntra_fashion_dataset$Ratings),
  mean(myntra_fashion_dataset$Ratings, na.rm = TRUE),
  myntra_fashion_dataset$Ratings)
myntra_fashion_dataset$Reviews <-
  ifelse(is.na(myntra_fashion_dataset$Reviews),
```

```
mean(myntra_fashion_dataset$Reviews, na.rm = TRUE),
myntra_fashion_dataset$Reviews)
```

Let's remove the unnecessary columns. (URL, SizeOption, Description)

```
myntra_fashion_dataset = subset(myntra_fashion_dataset, select = -
c(URL, SizeOption, Description) )
head(myntra_fashion_dataset)
```

	Product_id	BrandName	Category	Individual_category
1	2296012	Roadster	Bottom Wear	jeans
2	13780156	LOCOMOTIVE	Bottom Wear	track-pants
3	11895958	Roadster	Topwear	shirts
4	4335679	Zivame	Lingerie & Sleep Wear	shapewear
5	11690882	Roadster	Western	tshirts
6	2490950	Mast & Harbour	Western	tops
		category_by_Gender	DiscountPrice (in Rs)	OriginalPrice (in Rs)
		DiscountOffer		
1	Men		824	1499
3.002001				
2	Men		517	1149
4.786771				
3	Men		629	1399
3.931380				
4	Women		893	1295
2.393822				
5	Women		564	599
5.843072				
6	Women		559	599
6.677796				
	Ratings	Reviews		
1	3.9	999		
2	4.0	999		
3	4.3	999		
4	4.2	999		
5	4.2	999		
6	4.4	999		

```
missing_counts <- colSums(is.na(myntra_fashion_dataset)) # checking
for the missing value count
print(missing_counts)
```

	Product_id	BrandName	Category
	0	0	0
Individual_category		category_by_Gender	DiscountPrice (in Rs)
	0	0	0
OriginalPrice (in Rs)		DiscountOffer	Ratings
	0	0	0
	Reviews		
	0		

```
sum(is.na(myntra_fashion_dataset$'OriginalPrice (in Rs)'))
[1] 0
```

Here we can see that all the missing values has been handled. Now the data can be used for understanding statistical modelling techniques.

```
colnames(myntra_fashion_dataset)[6] <- c('DiscountPrice') # changing
the column names for convinience
colnames(myntra_fashion_dataset)[7] <- c('OriginalPrice')

missing_counts <- colSums(is.na(myntra_fashion_dataset))
print(missing_counts)



|                     | Product_id    | BrandName     | Category |
|---------------------|---------------|---------------|----------|
| Individual_category | 0             | 0             | 0        |
| category_by_Gender  | DiscountPrice | OriginalPrice |          |
| DiscountOffer       | 0             | 0             | 0        |
|                     | Ratings       | Reviews       |          |
|                     | 0             | 0             |          |



summary(myntra_fashion_dataset)



|                     | Product_id       | BrandName        | Category         |
|---------------------|------------------|------------------|------------------|
| Individual_category | Min. : 27399     | Length:526564    | Length:526564    |
|                     | Length:526564    |                  |                  |
|                     | 1st Qu.:13880530 | Class :character | Class :character |
|                     | Class :character |                  |                  |
|                     | Median :15971057 | Mode :character  | Mode :character  |
|                     | Mode :character  |                  |                  |
|                     | Mean :15069387   |                  |                  |
|                     | 3rd Qu.:17347414 |                  |                  |
|                     | Max. :18464352   |                  |                  |
| category_by_Gender  | Min. : 99        | Min. : 99        | Min. : 0.000     |
| Length:526564       | 1st Qu.: 736     | 1st Qu.: 1299    | 1st Qu.: 1.000   |
| Class :character    | Median : 1169    | Median : 1999    | Median : 2.155   |
| Mode :character     | Mean : 1507      | Mean : 2414      | Mean : 3.911     |
|                     | 3rd Qu.: 1890    | 3rd Qu.: 2899    | 3rd Qu.: 3.339   |
|                     | Max. :90000      | Max. :90000      | Max. :89.047     |
| Ratings             | Min. :1.000      | Min. : 0.00      |                  |


```

```

1st Qu.:4.095   1st Qu.: 40.00
Median :4.095   Median : 61.99
Mean   :4.095   Mean   : 61.99
3rd Qu.:4.095   3rd Qu.: 61.99
Max.    :5.000   Max.   :999.00

```

Converting categorical data to Numerical Data

We do have some categorical data in our dataset. Like the Gender category, this columnn only has Men and Women as its unique values. We can convert them to numerical data. Here we will be assigning Men as 0 and Women as 1.

```

gender_category <- unique(myntra_fashion_dataset$category_by_Gender)
gender_category <- length(gender_category)
gender_category

[1] 2

myntra_fashion_dataset <- myntra_fashion_dataset %>%
  mutate(category_by_Gender = if_else(category_by_Gender == "Men", 0,
  1))

```

Here, we have 8 unique values for the category variable. We can convert this into numerical data and use it for statistical analysis

```

unique_value_category <- unique(myntra_fashion_dataset$Category)
unique_no <- length(unique_value_category)
unique_no
unique_value_category

[1] 8

[1] "Bottom Wear"           "Topwear"
[3] "Lingerie & Sleep Wear" "Western"
[5] "Sports Wear"           "Indian Wear"
[7] "Plus Size"              "Inner Wear & Sleep Wear"

myntra_fashion_dataset$Category <-
as.numeric(factor(myntra_fashion_dataset$Category, levels = c('Bottom
Wear','Topwear','Lingerie & Sleep Wear','Western','Sports
Wear','Indian Wear','Plus Size','Inner Wear & Sleep Wear')))

head(myntra_fashion_dataset$Category) # here the 8 unique categories
has been converted to 1-8 numerical

[1] 1 1 2 3 4 4

```

Dataset division into training and test dataset

```
# Dividing our dataset into training and test set
set.seed(1111)
n = floor(0.8 * nrow(myntra_fashion_dataset)) #find the number
corresponding to 80% of the data
index = sample(seq_len(nrow(myntra_fashion_dataset)), size = n)
#randomly sample indices to be included in the training set

train = myntra_fashion_dataset[index, ] #set the training set to be
the randomly sampled rows of the dataframe
test = myntra_fashion_dataset[-index, ] #set the testing set to be the
remaining rows
cat("There are", dim(train)[1], "rows and", dim(train)[2], "columns in
the training set. ")
cat("There are", dim(test)[1], "rows and", dim(test)[2], "columns in the
testing set.")
```

There are 421251 rows and 10 columns in the training set. There are 105313 rows and 10 columns in the testing set.

```
str(train)

tibble [421,251 × 10] (S3: tbl_df/tbl/data.frame)
$ Product_id      : num [1:421251] 14737378 17702838 17225968
10182339 17422112 ...
$ BrandName       : chr [1:421251] "Mr Bowerbird" "ATTITUDE"
"Jinfo" "Style Quotient" ...
$ Category        : num [1:421251] 2 2 6 4 6 4 3 7 2 8 ...
$ Individual_category: chr [1:421251] "tshirts" "tshirts" "palazzos"
"tops" ...
$ category_by_Gender : num [1:421251] 0 0 1 1 1 1 1 1 0 0 ...
$ DiscountPrice   : num [1:421251] 499 779 1049 1535 1880 ...
$ OriginalPrice   : num [1:421251] 999 1299 2498 1599 3760 ...
$ DiscountOffer   : num [1:421251] 5.01 3.08 2.32 4 1.33 ...
$ Ratings         : num [1:421251] 4.7 4.09 4.09 4.2 4.09 ...
$ Reviews         : num [1:421251] 28 62 62 272 62 ...
```

Here, we are creating a subset of our training dataset and including just the numerical data

```
myntra_fashion_subset <- subset(train, select = -c(Product_id,
BrandName, Individual_category))
head(myntra_fashion_subset)
```

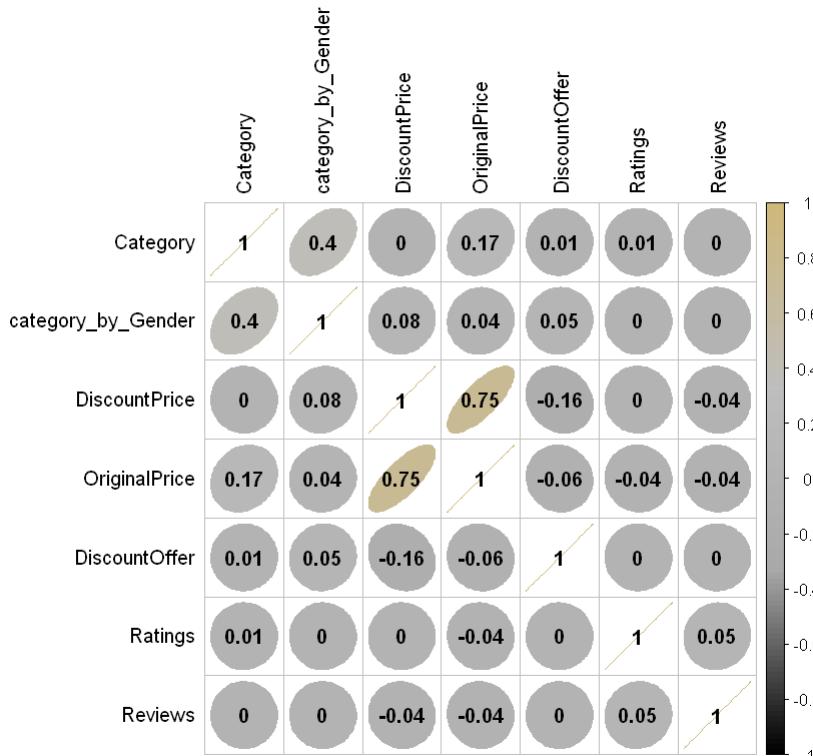
	Category	category_by_Gender	DiscountPrice	OriginalPrice
	DiscountOffer			
1	2	0	499	999
2	2	0	779	1299
3	6	1	1049	2498

4	4	1	1535	1599	4.002502
5	6	1	1880	3760	1.329787
6	4	1	949	999	5.005005
Ratings Reviews					
1	4.700000	28.00000			
2	4.094892	61.99082			
3	4.094892	61.99082			
4	4.200000	272.00000			
5	4.094892	61.99082			
6	4.600000	13.00000			

Correlation matrix

Let's create a correlation matrix to understand the dependency of each feature on one another.

```
col4 = colorRampPalette(c("black", "darkgrey", "grey", "#CFB87C"))
corrplot(cor(myntra_fashion_subset[]), method = "ellipse", col =
col4(100), addCoef.col = "black", tl.col = "black")
```



From the correlation matrix, we can say that there is a high correlation between the `OriginalPrice` and `DiscountPrice`

Part B: Regression Modelling

Let's perform linear regression on our training dataset with numerical features. We will be using `DiscountPrice` as our response variable and all other features as predictor variables

```
lm_myntra <- lm(DiscountPrice ~ ., data = myntra_fashion_subset)
summary(lm_myntra)

Call:
lm(formula = DiscountPrice ~ ., data = myntra_fashion_subset)

Residuals:
    Min      1Q  Median      3Q     Max 
-15135   -436    -39     408   42672 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 1.719e+02  1.761e+01   9.76  <2e-16 ***
Category    -1.253e+02  7.271e-01 -172.34  <2e-16 ***
category_by_Gender 3.491e+02  2.845e+00  122.71  <2e-16 ***
OriginalPrice 5.237e-01  6.634e-04  789.46  <2e-16 ***
DiscountOffer -1.577e+01  1.319e-01 -119.58  <2e-16 ***
Ratings       1.061e+02  4.221e+00   25.14  <2e-16 ***
Reviews       -1.748e-01  1.658e-02  -10.54  <2e-16 ***  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 810.5 on 421244 degrees of freedom
Multiple R-squared:  0.6106,    Adjusted R-squared:  0.6106 
F-statistic: 1.101e+05 on 6 and 421244 DF,  p-value: < 2.2e-16
```

`lm_myntra` is the linear regression model with `DiscountPrice` as the response variable and all other are the predictor variable.

From the summary output we can see that all the features from our dataset seem to be relevant for the model to be a good fit. Considering significance value as $\alpha=0.05$, the P-value for all the features seems to be less than the considered significance value. But let's not conclude now, we can go ahead and use our model selection techniques to understand which size of feature is the best fit for the model.

Part C: Backward Selection and Diagnostics of model

Let's perform backward selection process to see which size of features is the best fit.

Removing one feature from the model and analyze.

```
updated_lm_myntra = update(lm_myntra, . ~ . -Category)
summary(updated_lm_myntra)

Call:
lm(formula = DiscountPrice ~ category_by_Gender + OriginalPrice +
    DiscountOffer + Ratings + Reviews, data = myntra_fashion_subset)

Residuals:
    Min      1Q  Median      3Q     Max 
-14568   -415    -40     402   44217 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) -1.338e+02  1.813e+01  -7.38  1.58e-13 ***
category_by_Gender 1.550e+02  2.703e+00   57.33 < 2e-16 ***
OriginalPrice   5.042e-01  6.763e-04  745.52 < 2e-16 ***
DiscountOffer   -1.582e+01  1.365e-01 -115.87 < 2e-16 ***
Ratings         9.698e+01  4.367e+00   22.21 < 2e-16 ***
Reviews        -1.918e-01  1.716e-02  -11.18 < 2e-16 ***
...
Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 838.6 on 421245 degrees of freedom
Multiple R-squared:  0.5832,    Adjusted R-squared:  0.5832 
F-statistic: 1.179e+05 on 5 and 421245 DF,  p-value: < 2.2e-16
```

Here we removed the Category feature to see if this model is a best fit or not. The update model still says all the features seems good. Using a significance value as $\alpha=0.05$, the P-value for all the features seems to be less than the considered significance value.

```
updated_lm_myntra_2 = update(updated_lm_myntra, . ~ . -
category_by_Gender)
summary(updated_lm_myntra_2)

Call:
lm(formula = DiscountPrice ~ OriginalPrice + DiscountOffer +
    Ratings + Reviews, data = myntra_fashion_subset)

Residuals:
    Min      1Q  Median      3Q     Max 
          Min      1Q  Median      3Q     Max 

```

```
-14571 -428 -63 416 44133
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)		
(Intercept)	-4.063e+01	1.813e+01	-2.242	0.025 *		
OriginalPrice	5.058e-01	6.783e-04	745.644	<2e-16 ***		
DiscountOffer	-1.541e+01	1.368e-01	-112.617	<2e-16 ***		
Ratings	9.722e+01	4.384e+00	22.178	<2e-16 ***		
Reviews	-1.909e-01	1.723e-02	-11.080	<2e-16 ***		

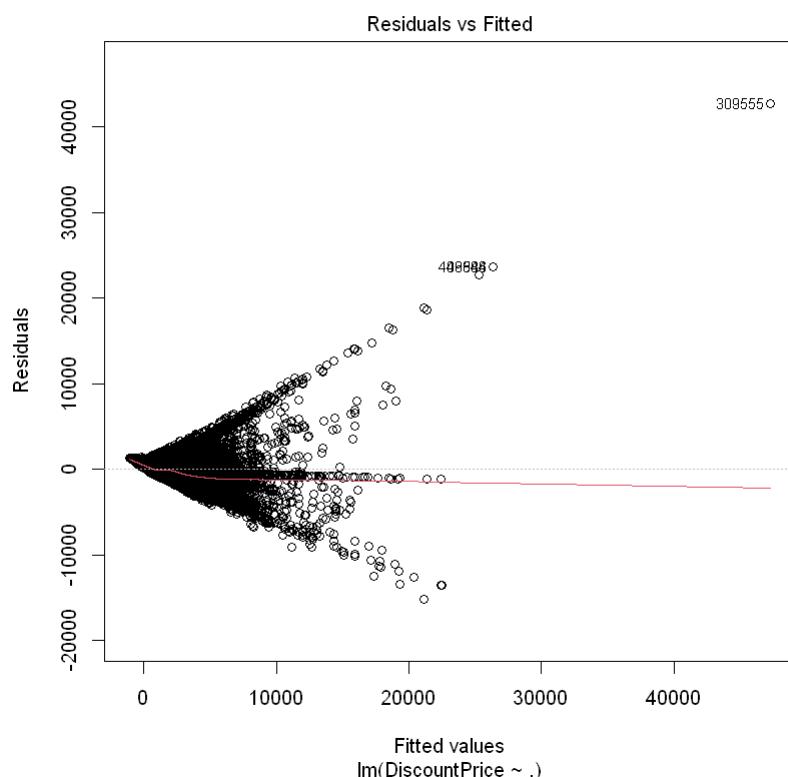
Signif. codes:	0 '***'	0.001 '**'	0.01 '*'	0.05 '.'	0.1 ' '	1

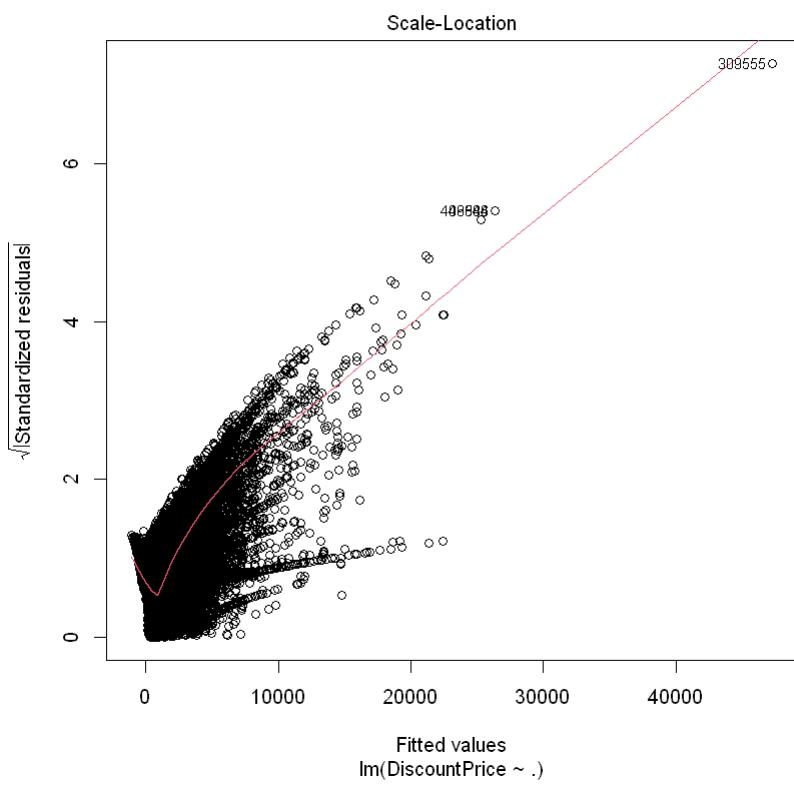
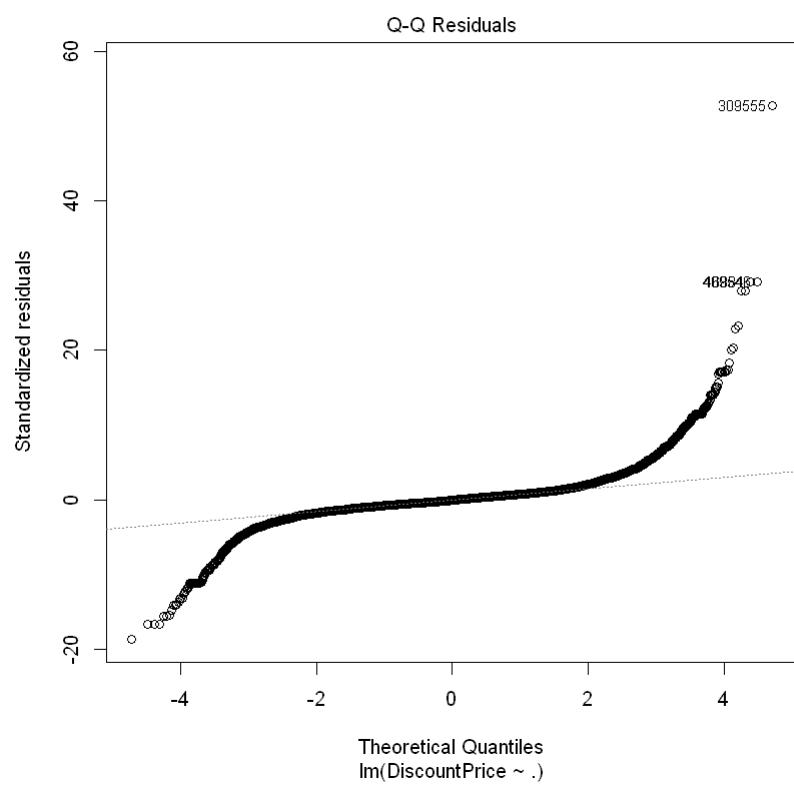
Residual standard error: 841.8 on 421246 degrees of freedom
Multiple R-squared: 0.5799, Adjusted R-squared: 0.5799
F-statistic: 1.454e+05 on 4 and 421246 DF, p-value: < 2.2e-16

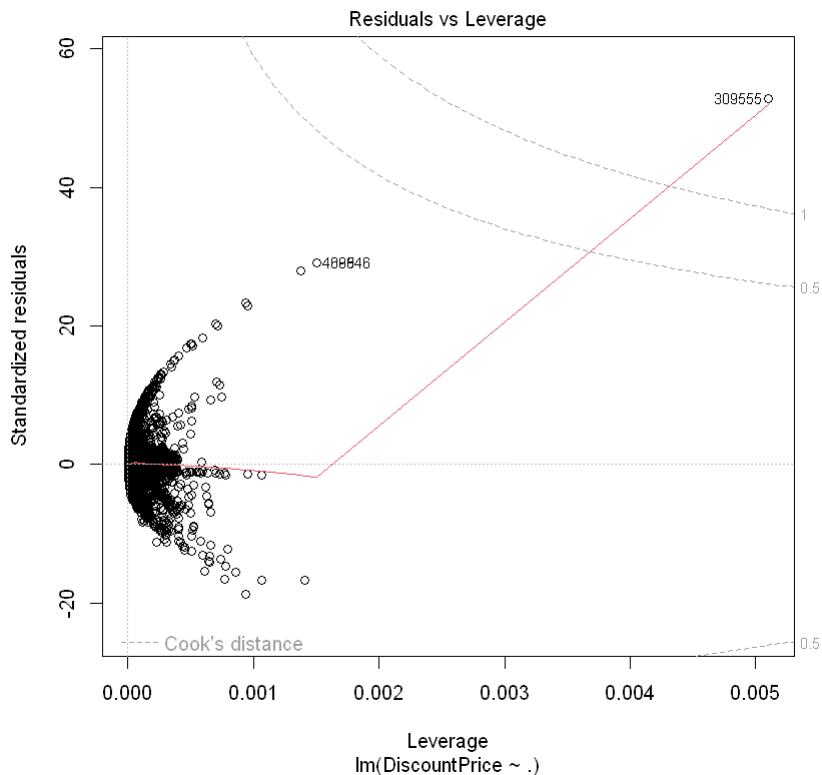
I went ahead with updating our model by removing the `category_by_gender` feature. Although after removing certain features, we can see that all the features still seem to be the best fits for a model

Now Let's plot the diagnostics plot for our model

```
plot(lm_myntra)
```







- Residual Plot: It is difficult to say definitively whether the assumptions of the linear model have been met. There appears to be a slight trend in the residuals, with positive residuals at lower fitted values and negative residuals at higher fitted values. This could indicate that the linear model is not capturing the relationship between the discount price and the feature variables very well.
- Q-Q Plot: In the Q-Q plot, the points deviate from the straight line, particularly for the lower and upper quantiles of the residuals. This suggests that the residuals are not normally distributed. The data is very skewed and has a lot of outliers. There appears to be a heavier tail on both ends of the distribution than would be expected in a normal distribution. This could indicate that the linear model is not capturing the relationship between the discount price and the feature variables very well.
- Scale-Location PPlot: In this plot, it appears that the variance of the residuals might be increasing slightly with increasing fitted values. This suggests that the linear model might not have constant variance.
- Residual Vs Leverage Plot: there appears to be a slight trend in the residuals, with positive residuals at lower fitted values and negative residuals at higher fitted values.

```
mynt_fashion_subset_test <- subset(test, select = -c(Product_id,
BrandName, Individual_category))
head(mynt_fashion_subset_test)
```

	Category	category_by_Gender	DiscountPrice	OriginalPrice	
DiscountOffer	Ratings	824	1499	3.002001	
1	1	0			

3.9				
2 4	1	564	599	5.843072
4.2				
3 1	0	2749	2749	0.000000
3.5				
4 6	1	696	1699	3.472631
4.2				
5 4	1	1449	1499	3.335557
4.2				
6 4	1	1548	1598	3.128911
4.2				
Reviews				
1	999			
2	999			
3	996			
4	995			
5	993			
6	990			

Calculating MSPE values for each model

```
calculate_MSPE <- function(model, test_data) {
  predictions <- predict(model, newdata = test_data)
  return(mean((test_data$DiscountPrice - predictions)^2))
}

temp_mspe <- calculate_MSPE(updated_lm_myntra,
myntra_fashion_subset_test) # MSPE for model with 5 features
temp_mspe

[1] 709491.1

temp_mspe <- calculate_MSPE(updated_lm_myntra_2,
myntra_fashion_subset_test) # MSPE for model with 4 features
temp_mspe

[1] 715233.6

temp_mspe <- calculate_MSPE(lm_myntra, myntra_fashion_subset_test) # MSPE for model with 6 features
temp_mspe

[1] 663293.5
```

- 715233.647652179 This is the MSPE for model with 4 features
- 663293.548101405 This is the MSPE for model with 6 features
- 709491.144712724 This is the MSPE for model with 5 features

we can see that all the MSPE values are really high, but one possible reason for this high MSPE could be the huge dataset of approx 5 lakh rows. Apart from that we can say the the model with 6 feature variables give the best model. The model with 6 feature variables has comparatively less MSPE value than the other model's MSPE value.

Transforming the data

Here we are transforming our dataset. We are applying the square root transformation.

```

myntra_fashion_subset$Category_transformed <-
  sqrt(myntra_fashion_subset$Category)
myntra_fashion_subset$category_by_Gender_transformed <-
  sqrt(myntra_fashion_subset$category_by_Gender)
myntra_fashion_subset$DiscountPrice_transformed <-
  sqrt(myntra_fashion_subset$DiscountPrice)
myntra_fashion_subset$OriginalPrice_transformed <-
  sqrt(myntra_fashion_subset$OriginalPrice)
myntra_fashion_subset$DiscountOffer_transformed <-
  sqrt(myntra_fashion_subset$DiscountOffer)
myntra_fashion_subset$Ratings_transformed <-
  sqrt(myntra_fashion_subset$Ratings)
myntra_fashion_subset$Reviews_transformed <-
  sqrt(myntra_fashion_subset$Reviews)

lm_myntra_transformed <- lm(DiscountPrice_transformed ~
  Reviews_transformed +
  Ratings_transformed + DiscountOffer_transformed + OriginalPrice_transformed +
  category_by_Gender_transformed + Category_transformed, data =
  myntra_fashion_subset)
summary(lm_myntra_transformed)

Call:
lm(formula = DiscountPrice_transformed ~ Reviews_transformed +
    Ratings_transformed + DiscountOffer_transformed +
    OriginalPrice_transformed +
    category_by_Gender_transformed + Category_transformed, data =
    myntra_fashion_subset)

Residuals:
    Min      1Q  Median      3Q     Max 
-53.850 -5.260 -0.038  5.777 107.053 

Coefficients:
                                         Estimate Std. Error t value Pr(>|t|)    
(Intercept)                         11.5441030   0.3170438   36.41   <2e-16  

```

```

***  

Reviews_transformed           -0.0550292  0.0037910  -14.52  <2e-16  

***  

Ratings_transformed           3.6991058  0.1540423   24.01  <2e-16  

***  

DiscountOffer_transformed     -2.5312848  0.0096502 -262.30  <2e-16  

***  

OriginalPrice_transformed    0.6046812  0.0007859  769.39  <2e-16  

***  

category_by_Gender_transformed 4.6556745  0.0286674  162.40  <2e-16  

***  

Category_transformed          -4.7804624  0.0274135 -174.38  <2e-16  

***  

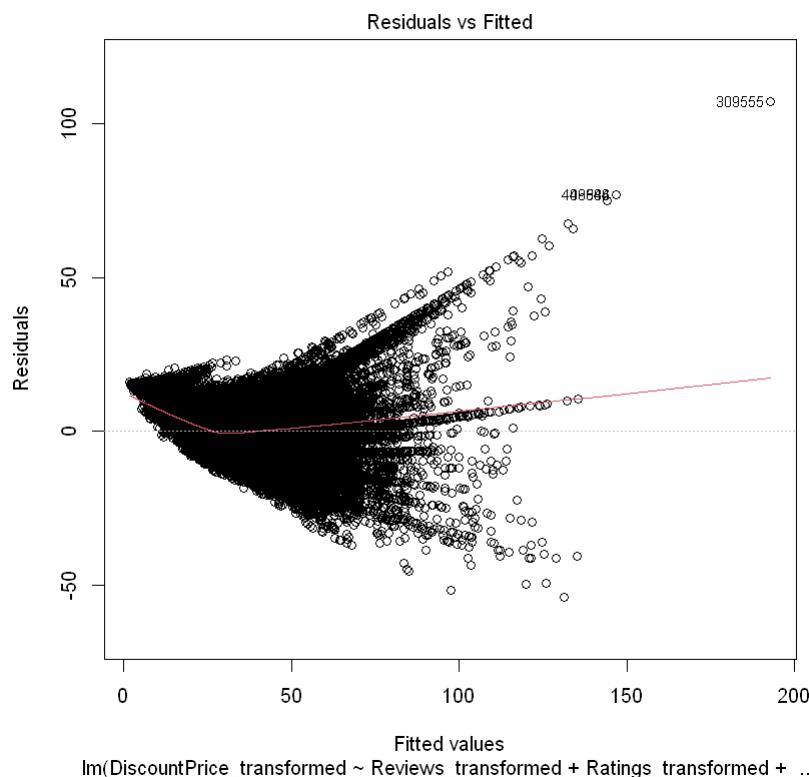
---  

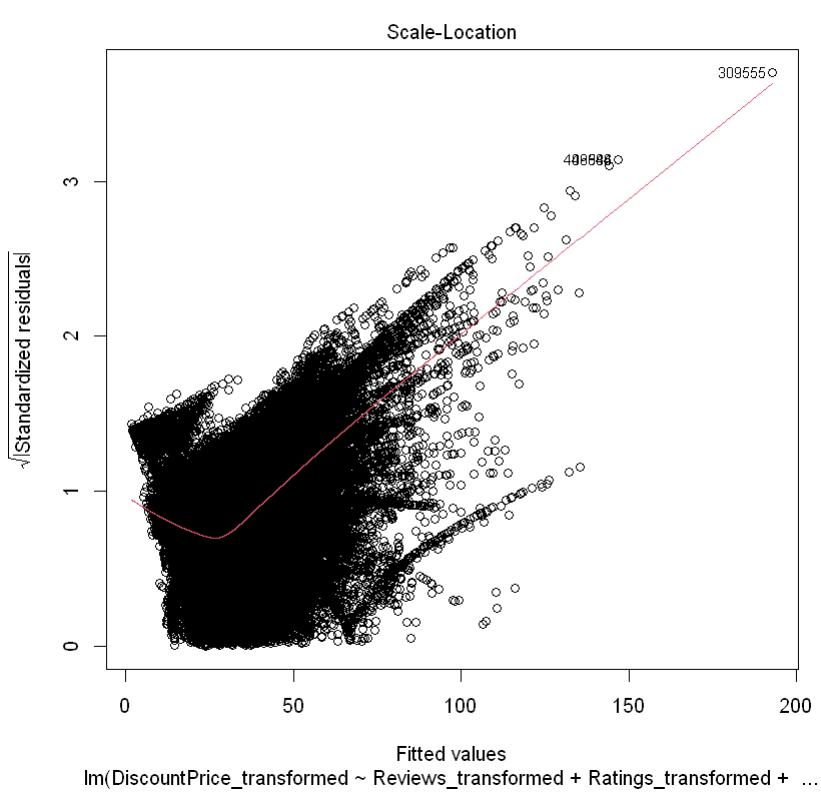
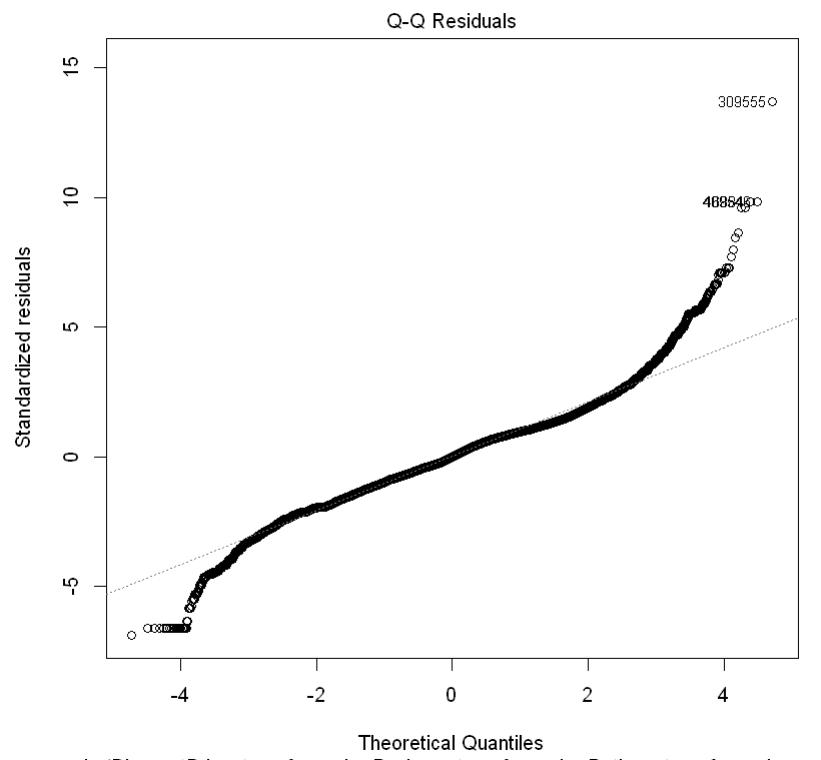
Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

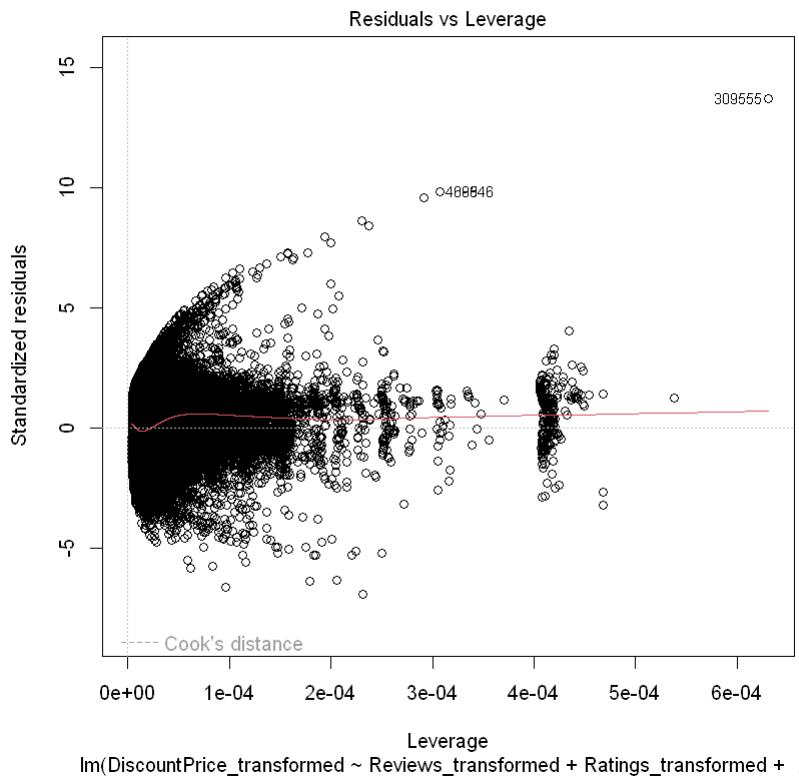
```

Residual standard error: 7.815 on 421244 degrees of freedom
Multiple R-squared: 0.6442, Adjusted R-squared: 0.6442
F-statistic: 1.271e+05 on 6 and 421244 DF, p-value: < 2.2e-16

```
plot(lm_myntra_transformed) # diagnostic plots
```







- Residual Plot: the residuals appear to be scattered somewhat randomly around the horizontal line at $y=0$. There is a slight trend where the residuals seem to be more positive for higher fitted values, but it is not a strong trend. This suggests that the linear model might be doing a decent job of capturing the relationship between the transformed discount price and the feature variables. However, there might still be a slight violation of the assumption of homoscedasticity.
- Q-Q Plot: In the Q-Q plot, the points deviate slightly from the straight line, particularly for the tails of the distribution. This suggests that the residuals are not perfectly normal. However, the deviation is not severe as compared to the un-transformed data, and it can be acceptable for linear regression.
- Scale-Location PPlot: In this plot, the residuals appear to be scattered somewhat randomly around the red line. There is a slight trend where the residuals seem to have a larger spread for higher fitted values, but it is not a strong trend. This suggests that the linear model might not have perfect homoscedasticity, but the violation might be mild.
- Residual Vs Leverage Plot: the standardized residuals appear to be scattered somewhat randomly around the horizontal line at $y=0$. This suggests that the linear model might be doing a decent job of capturing the relationship between the transformed discount price and the feature variables. However, there might still be a slight violation of the assumption of homoscedasticity.

Part D: Forward Selection and Model selection

```
reg1 = regsubsets(DiscountPrice_transformed ~ Reviews_transformed + Ratings_transformed+DiscountOffer_transformed+OriginalPrice_transformed+category_by_Gender_transformed+Category_transformed, data = myntra_fashion_subset)
rs = summary(reg1)
rs$which
```

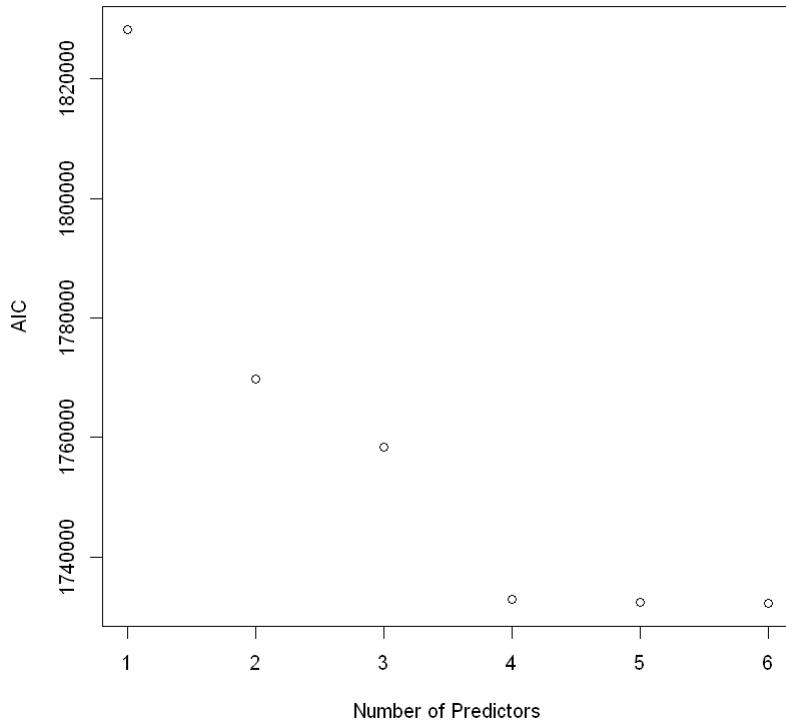
	(Intercept)	Reviews_transformed	Ratings_transformed	DiscountOffer_transformed
1	TRUE	FALSE	FALSE	FALSE
2	TRUE	FALSE	FALSE	TRUE
3	TRUE	FALSE	FALSE	TRUE
4	TRUE	FALSE	FALSE	TRUE
5	TRUE	FALSE	TRUE	TRUE
6	TRUE	TRUE	TRUE	TRUE

	OriginalPrice_transformed	category_by_Gender_transformed	Category_transformed
1	TRUE	FALSE	FALSE
2	TRUE	FALSE	FALSE
3	TRUE	FALSE	TRUE
4	TRUE	TRUE	TRUE
5	TRUE	TRUE	TRUE
6	TRUE	TRUE	TRUE

The above table shows that for each size, what features will give the best model.

To choose the best size, let use the model selection criterion i.e., AIC, BIC, MSPE, R^2 , R_a^2

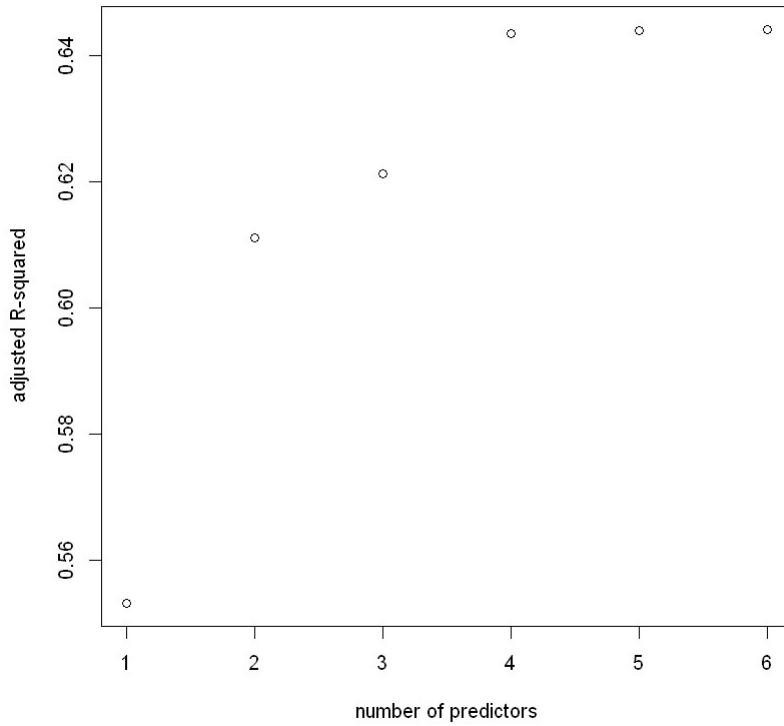
```
AIC = 2*(2:7) + n*log(rs$rss/n)
plot(AIC ~ I(1:6), xlab = "Number of Predictors", ylab = "AIC")
```



In this plot, we see that the model of size $k=6$ has the lowest AIC. That means that our model selection procedure has chosen:

```
discountprice = $\widehat{\beta}_0 + \widehat{\beta}_1 \times \$Category + \widehat{\beta}_2 \times \$category_by_Gender + \widehat{\beta}_3 \times \$OriginalPrice + \widehat{\beta}_4 \times \$DiscountOffer + \hat{\beta}_5 \times \$Ratings + \widehat{\beta}_6 \times \$Reviews.
```

```
plot(1:6, rs$adjr2, xlab = "number of predictors", ylab = "adjusted R-squared")
```



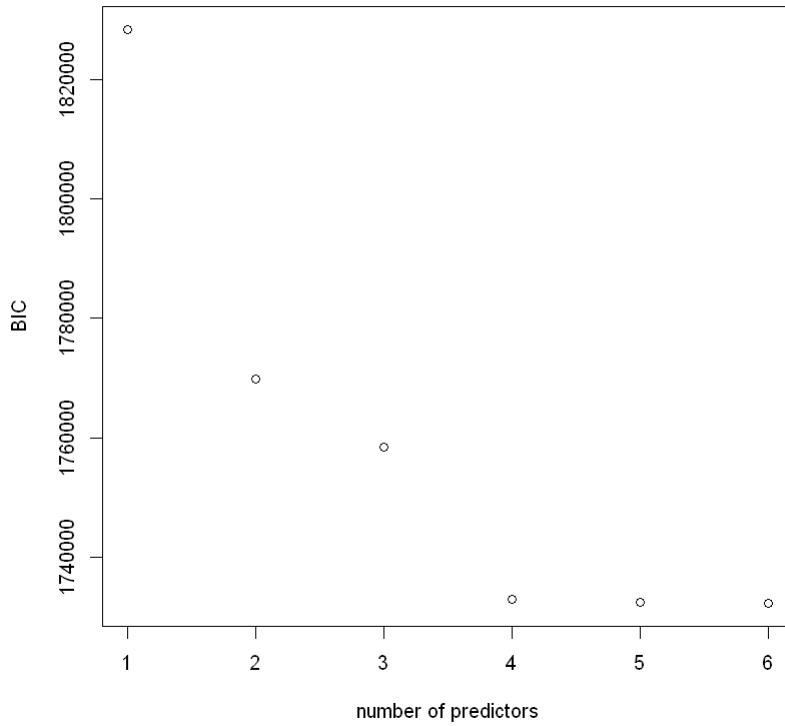
In this plot, we see that the model of size $k=6$ has the highest R_a^2 . That means that our model selection procedure has chosen:

```
discountprice = $\widehat{\beta}_0 + \widehat{\beta}_1 \times \$Category + \widehat{\beta}_2 \times \$category_by_Gender + \widehat{\beta}_3 \times \$OriginalPrice + \widehat{\beta}_4 \times \$DiscountOffer + \hat{\beta}_5 \times \$Ratings + \widehat{\beta}_6 \times \$Reviews.
```

Using the BIC formula as given below:

$$BIC(g(x; \hat{\beta})) = (p+1)\log(n) - 2\log L(\hat{\beta}),$$

```
BIC = log(n)*(2:7) + n*log(rs$rss/n)
plot(BIC ~ I(1:6), xlab = "number of predictors", ylab = "BIC")
```



In this plot, we see that the model of size $k=6$ has the lowest BIC. That means that our model selection procedure has chosen:

```
discountprice = $\widehat{\beta}_0 + \widehat{\beta}_1 \times \$Category + \widehat{\beta}_2 \times \$category_by_Gender + \widehat{\beta}_3 \times \$OriginalPrice + \widehat{\beta}_4 \times \$DiscountOffer + \hat{\beta}_5 \times \$Ratings + \widehat{\beta}_6 \times \$Reviews.
```

All the criterion give the value of $k=6$. This tells us the original model with all the numerical features gives the best model for linear regression

Part E: Hypothesis Testing and ANOVA:

Let's define some hypothesis

- **Null Hypothesis:** The Products with more reviews has more discounts.
- **Alternate Hypothesis:** There is no dependency of the reviews on the discounts

Now let's perform ANOVA to test these hypothesis.

```
anova_result <- aov(DiscountPrice ~ Reviews, data =
mynter_fashion_subset)
print(anova_result)
```

```

Call:
  aov(formula = DiscountPrice ~ Reviews, data =
myntra_fashion_subset)

Terms:
          Reviews   Residuals
Sum of Squares 961756987 709683282660
Deg. of Freedom           1      421249

Residual standard error: 1297.965
Estimated effects may be unbalanced

```

Looking at the ANOVA results from above:

The large Sum of Squares (961756987) for Reviews compared to Residuals (709683282660) says that the number of reviews explains a substantial portion of the variance in the square root of discount prices.

Additionally, let's look at the full ANOVA table for our transformed model

```

anova_result <- anova(lm_myntra_transformed)
print(anova_result)

Analysis of Variance Table

Response: DiscountPrice_transformed
                         Df  Sum Sq Mean Sq F value    Pr(>F)
Reviews_transformed       1    11759   11759 1.9254e+02 < 2e-16 ***
Ratings_transformed      1      219     219  3.5895e+00 0.05815 .
DiscountOffer_transformed 1   9230034  9230034 1.5113e+05 < 2e-16 ***
OriginalPrice_transformed 1  34988927 34988927 5.7290e+05 < 2e-16 ***
category_by_Gender_transformed 1   500008   500008 8.1870e+03 < 2e-16 ***
Category_transformed      1   1857214  1857214 3.0410e+04 < 2e-16 ***
Residuals                  421244 25726779            61
---
Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

From the 1st anova results, we saw that number of reviews explains a substantial portion of the variance in the square root of discount prices.

Now the full table gives us the P-values for all the features. All the features have a statistically significant effect on the transformed discount price at a significance level of alpha = 0.001. This

is indicated by the highly significant p-values ($\text{Pr}(>F)$) much lower than 0.001, we can also see looking at the (**),

that implies the features with the (**) indicate statistical significance.

We can, therefore say, **we failed to reject Null Hypothesis**, as there is significant evidence supporting the Null Hypothesis.

Part F: T-test

```
t_test_category <- t.test(myntra_fashion_dataset$DiscountOffer,  
myntra_fashion_dataset$category_by_Gender, data =  
myntra_fashion_dataset, subset = category_by_Gender %in% c("Women",  
"Men"))  
print(t_test_category)
```

```
Welch Two Sample t-test
```

```
data: myntra_fashion_dataset$DiscountOffer and  
myntra_fashion_dataset$category_by_Gender  
t = 248.9, df = 529230, p-value < 2.2e-16  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
 3.241426 3.292880  
sample estimates:  
mean of x mean of y  
3.9113011 0.6441477
```

The t-statistic ($t = 248.9$) is a very large positive value, indicating a substantial difference in the means between the two categories (women's and men's). The p-value is much smaller than 0.05, providing very strong evidence against the null hypothesis of no difference in means. There's a very high chance that this observed difference is not due to random chance.

The 95% confidence interval (3.241426, 3.292880) shows the range of values likely to contain the true difference in average discount between the categories. Since the entire interval is positive, we can be confident that the average discount for one group is higher than the other.

The average discount offered on products in the "Women" and "Men" categories is significantly different. This suggests that there might be generally higher discounts offered on the platform compared to specifically women's or men's categories.

Report:

Title: Myntra Fashion Store: Data Analysis

1. Introduction/Background

I have always been interested in Fashion and its changing trends over the time. The fashion industry needs thorough understanding of consumer behavior and preferences due to its fluctuating consumer trends and significant seasonal fluctuations. This study focuses on examining patterns in discount offers, product popularity, and customer reviews using data from Myntra, a well-known online apparel shop. These kinds of findings may be useful in customizing marketing plans and enhancing product offerings.

Nowadays, online shopping is the trend. People find it more convenient and time saving. Majority of the youth is dependent on the online fashion shopping. It comes with great perks like easy on budget, time saving, easy returns, etc. Online fashion retail is highly competitive, with success heavily dependent on understanding and reacting to fast-changing consumer preferences. Analyzing data from platforms like Myntra can provide insights into what influences purchase decisions, such as price, discounts, and product ratings.

This dataset is obtained from kaggle. Reason for collecting this dataset to understand the fashion trends and analyze the prices in the fashion industry. For this analysis, the data is taken from one of the leading online fashion store MYNTRA - (<https://www.myntra.com>) as their data source. This is an observational study.

Let's look at some of the questions that have caught my attention on looking at the dataset. (The 3rd question always makes me curious to think about the reason.)

- Does the product reviews affect the discounts?
- How does the price range affect customer reviews and ratings?
- Is there difference between the discounts offered based on Gender?

Some of the prior work that has been done with this dataset:

- <https://anshikanishad02030.medium.com/exploring-fashion-trends-a-comprehensive-analysis-of-myntra-fashion-dataset-2ecf9af4d664>
- https://github.com/aman9650/Myntra_Fashion_Clothing-EDA-Project

We can find more such projects online

2. Methods/Results (experimental design and data collection)

The data is taken from kaggle. I directly downloaded the dataset. The 'Myntra fasion dataset' is a huge database of approx. 526564 rows of data, with 13 columns. These columns include 'URL', 'Product_id', 'BrandName', 'Category', 'Individual_category', 'category_by_Gender', 'Description', 'DiscountPrice..in.Rs.', 'OriginalPrice..in.Rs.', 'DiscountOffer', 'SizeOption', 'Ratings' and 'Reviews'.

The dataset had many null values ad required a lot of data cleaning and data preprocessing for some columns like 'DiscountPrice..in.Rs.', 'DiscountOffer' and 'Ratings' and 'Reviews'. The data cleaning and data preprocessing is perfomed in the **Part A** of this project.

For exploratory data analysis, a correlation matrix is plotted, where we can say few features are highly correlated.(**Part A**).

For the statistical modelling, I have applied few methods and tried to answer the questions of interest.

- **Part B Regression Modelling:** This part includes a linear regression model on the obtained numerical dataset. This model seems to a good fit for now. But for further analysis, I included Backward selection on the trained model.
- **Part C Backward Selection and Diagnostics of model:** This part includes the Backward selection on the trained model and the diagnostic plots to understand the statistical assumption. It was difficult to make a point on the assumption of linear model. Therefore, applied a squareroot transformation on the dataset and re-run the model. Although there is a slight violation of homoscedasticity and there are few outliers in the dataset, yet this transformed model seems a good fit. Here, we also applied the MSPE criterion to find the best fit model.
- **Part D Forward Selection and Model selection:** This part includes the Forward selection, here using the different criterion, namely AIC, BIC, R^2 , R_a^2 , I was able to choose the best number of features to get the good fit. $K=6$ gives the best model for all the criterion
- **Part E Hypothesis Testing and ANOVA:** This part includes a Hypothesis Testing, the hypothesis states, whether there is a dependency on discounts and review, to test this hypothesis, I used ANOVA, which clearly indicates that the there is a dependency on discounts and review.
- **Part F T-test:** This part includes a t-test to see whether the discounts vary based on gender. Using the t-test results, it can be concluded that discounts seem to vary based on gender.

These methods have been specifically used to answer the questions of interest and to understand which features truly help determining the discountprice on a product.

3. Conclusions

This project helped me to understand how some questions in the real world can simply be answered using data and performing statistical analysis on that data. Statistical analysis provides a concrete reason on the statements/question(which are usually called hypothesis) we get. To highlight on the fashion database analysis, we got to know that the reviews affect the discount on a product. There is a difference in the discounts offered to Men and women.

I want to extend my research on understanding the sales in the fashion industry, does the review and discounts affect the sale of a product. I also want to build a personalized outfit

recommendation system or a dynamic price optimizer that would optimize pricing strategies dynamically based on demand, inventory levels, and trends.